
Chapter 3 Relational Model

3. Relational Model

3.1 Structure of a relational database

3.2 The relational algebra

3.2.1 Select, Project, Product, Union, Difference, Rename, Intersection, Division, Assignment, Natural Join, Outer Join, Aggregate functions, Generalized projection.

3.2.2 Database manipulation: Insertion, Update, and Deletion

Structure of Relational Databases

- A relational database consists of collection of tables, each of which is assigned a unique name.
- A row in a table represents a relationship among a set of values.
- A table is a collection of relationships.
- The concept of table closely corresponds to mathematical concept of relations.
- Thus, the relational model represents the database of collection of relations.

Basic Structure (1)

- For each attributes, there is a set of permitted values, called the **domain** of that attribute.
- E.g.: for branch-name attribute, the domain is the set of all branch names.
- If account has D1 that is set of all account numbers, D2 is set of all branch-name, and D3 the set of all balance, then any row of account must consist of a 3-tuple(v_1, v_2, v_3),
 - Where, v_1 is an account number that is in domain D1
 - v_2 is a branch-name that is in domain D2
 - v_3 is a balance that is in domain D3
- Therefore: account is a subset of all possible rows. Thus, account is a subset of
$$D1 \times D2 \times D3$$
- In general a table of n attributes must be a subset of
$$D_1 \times D_2 \times D_3 \times \dots \times D_{n-1} \times D_n$$

Basic Structure (2)

- Formally, given sets D_1, D_2, \dots, D_n domains, a relation r is a subset of $D_1 \times D_2 \times \dots \times D_n$
- *In other words, Mathematically, a relation is a subset of cartesian product of a list of domains.*
- *Thus a relation is a set of n -tuples (a_1, a_2, \dots, a_n) , $a_i \in D_i$*

Example: if

customer-name = {Jones, Smith, Curry, Lindsay}

customer-street = {Main, North, Park}

customer-city = {Munich, Stuttgart, Berlin}

Then $r = \{(Jones, Main, Munich),$
 $(Smith, North, Stuttgart),$
 $(Curry, North, Stuttgart),$
 $(Lindsay, Park, Berlin)\}$

is a relation over customer-name \times customer-street \times customer-city.

Attribute Types

- Each attribute of a relation has a name.
- The set of allowed values for each attribute is called the domain of that attribute.
- Attribute values are (normally) required to be atomic, that is, indivisible.
 - E.g. Multivalued attributes values are not atomic
 - E.g. Composite attributes values are not atomic
 - The special value null is a member of every domain.

Relational Schema

- A_1, A_2, \dots, A_n are attributes
- $R=(A_1, A_2, \dots, A_n)$ is a relation schema
 - E.g.: Customer-schema=(customer-name, customer street, customer city)
- $r(R)$ is a relation on relation schema R
 - E.g.: customer(Customer-Schema)
 - means *customer* is a relation on Customer-schema by *customer(Customer-schema)*

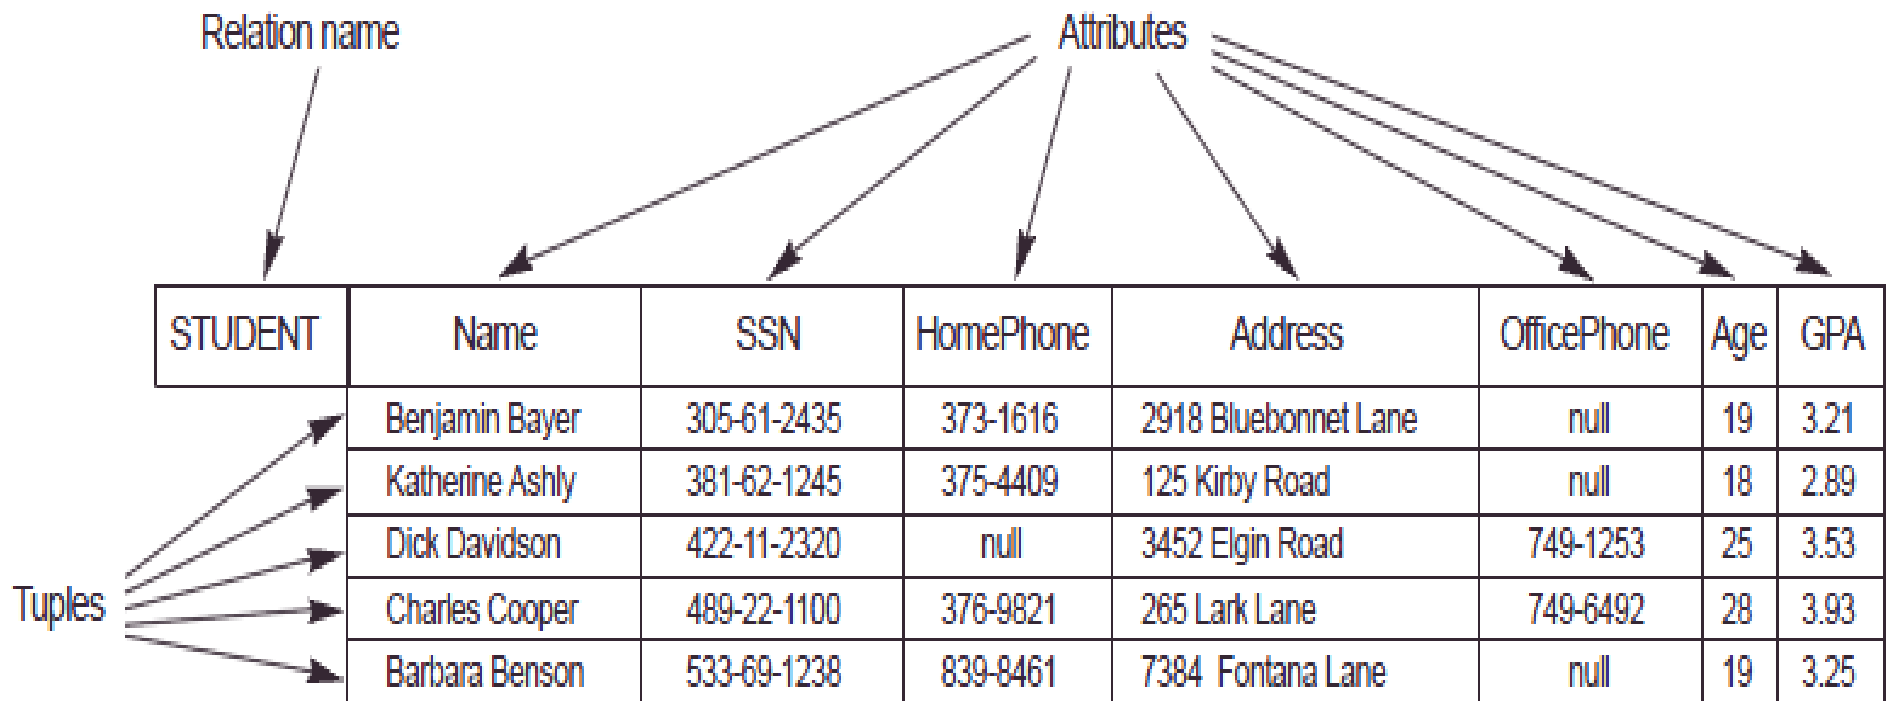
Naming convention

- Relation schema names begin with uppercase
- Relation names are in lowercase

Relation Instance

- The current values of a relation are specified by a table. Thus, a relation state at a given time is called a relation instance.
- An element t of r is a tuple, represent by row in a table.

Attributes and Tuples of a relation Student



Relation student is

student(Name, SSN, HomePhone, Address, OfficePhone, Age, GPA)

Relations are Unordered

- Order of tuples is irrelevant (tuples may be stored in an arbitrary order)
- E.g. account relation with unordered tuples

| <i>account-number</i> | <i>branch-name</i> | <i>balance</i> |
|-----------------------|--------------------|----------------|
| A-101 | Downtown | 500 |
| A-215 | Mianus | 700 |
| A-102 | Perryridge | 400 |
| A-305 | Round Hill | 350 |
| A-201 | Brighton | 900 |
| A-222 | Redwood | 700 |
| A-217 | Brighton | 750 |

Database

- A database consists of multiple relations.
- Information about an enterprise is broken up into parts, with each relation storing one part of the information

- E.g.:

| | |
|-------------------|--|
| <i>Account:</i> | stores information about accounts |
| <i>Depositor:</i> | stores information about which customer owns which account |
| <i>Customer:</i> | stores information about customers |

- Storing all information as a single relation such as
bank(account-number, balance, customer-name, . . .)
Results in
 - ❖ repetition of information (e.g. two customers own an account)
 - ❖ the need for null values (e.g. represent a customer without an account)

The *customer* relation

| <i>customer-name</i> | <i>customer-street</i> | <i>customer-city</i> |
|----------------------|------------------------|----------------------|
| Adams | Spring | Pittsfield |
| Brooks | Senator | Brooklyn |
| Curry | North | Rye |
| Glenn | Sand Hill | Woodside |
| Green | Walnut | Stamford |
| Hayes | Main | Harrison |
| Johnson | Alma | Palo Alto |
| Jones | Main | Harrison |
| Lindsay | Park | Pittsfield |
| Smith | North | Rye |
| Turner | Putnam | Stamford |
| Williams | Nassau | Princeton |

The *depositor* relation

| <i>customer-name</i> | <i>account-number</i> |
|----------------------|-----------------------|
| Hayes | A-102 |
| Johnson | A-101 |
| Johnson | A-201 |
| Jones | A-217 |
| Lindsay | A-222 |
| Smith | A-215 |
| Turner | A-305 |

Keys (1)

- Let R be a relation schema.
- *If a subset K of R is a superkey for R , then in relation $r(R)$, no two distinct tuples have the same values on all attributes in K . That is, if t_1 and t_2 are in r and $t_1 \neq t_2$, then $t_1[K] \neq t_2[K]$*
- *Example: {branch-name, branch-city}*

{branch-name}

Are both superkeys of branch, if no two branches can possibly have the same name.

- K is a candidate key if K is minimal.
- *Example: branch-name is a candidate key for branch, since it is a superkey (assuming no two branch can possibly have the same name), and no subset of it is a superkey*

Determining Keys from E-R sets

- Strong entity set: The primary key of the entity set becomes the primary key of the relation.
- Weak entity set: The primary key of the relation consists of union of the primary key of the strong entity set and the discriminator of the weak entity set.
- Relationship set: The union of primary keys of the related entity sets becomes a super key of the relation.
 - ❖ For binary many-to-one relationship sets, the primary key of “many” entity set becomes the relation’s primary key.
 - ❖ For one-to-one relationship sets, the relation’s primary key can be that of either entity set.
 - ❖ For many-to-many relationship sets, the union of primary keys of entity sets becomes the relation’s primary key.

Determining Keys from E-R sets

- Multivalued attributes:

- The primary key of entity set or relationship set, together with the attribute C becomes the relation's primary key. Where C is a column which holds an individual value of multivalued attribute M.

- Combined attributes:

- Binary many-to-one relationship set from entity set A to entity set B can be represented by a table consisting of the attribute of A and attributes of the relationship set. The primary key of the “many” entity set A becomes the primary key of the relation.
- Binary one-to-one relationship set from entity set A to entity set B can be constructed like in many-to-one. But either entity set's primary key is chosen as the primary key of the relation because both keys are candidate keys.

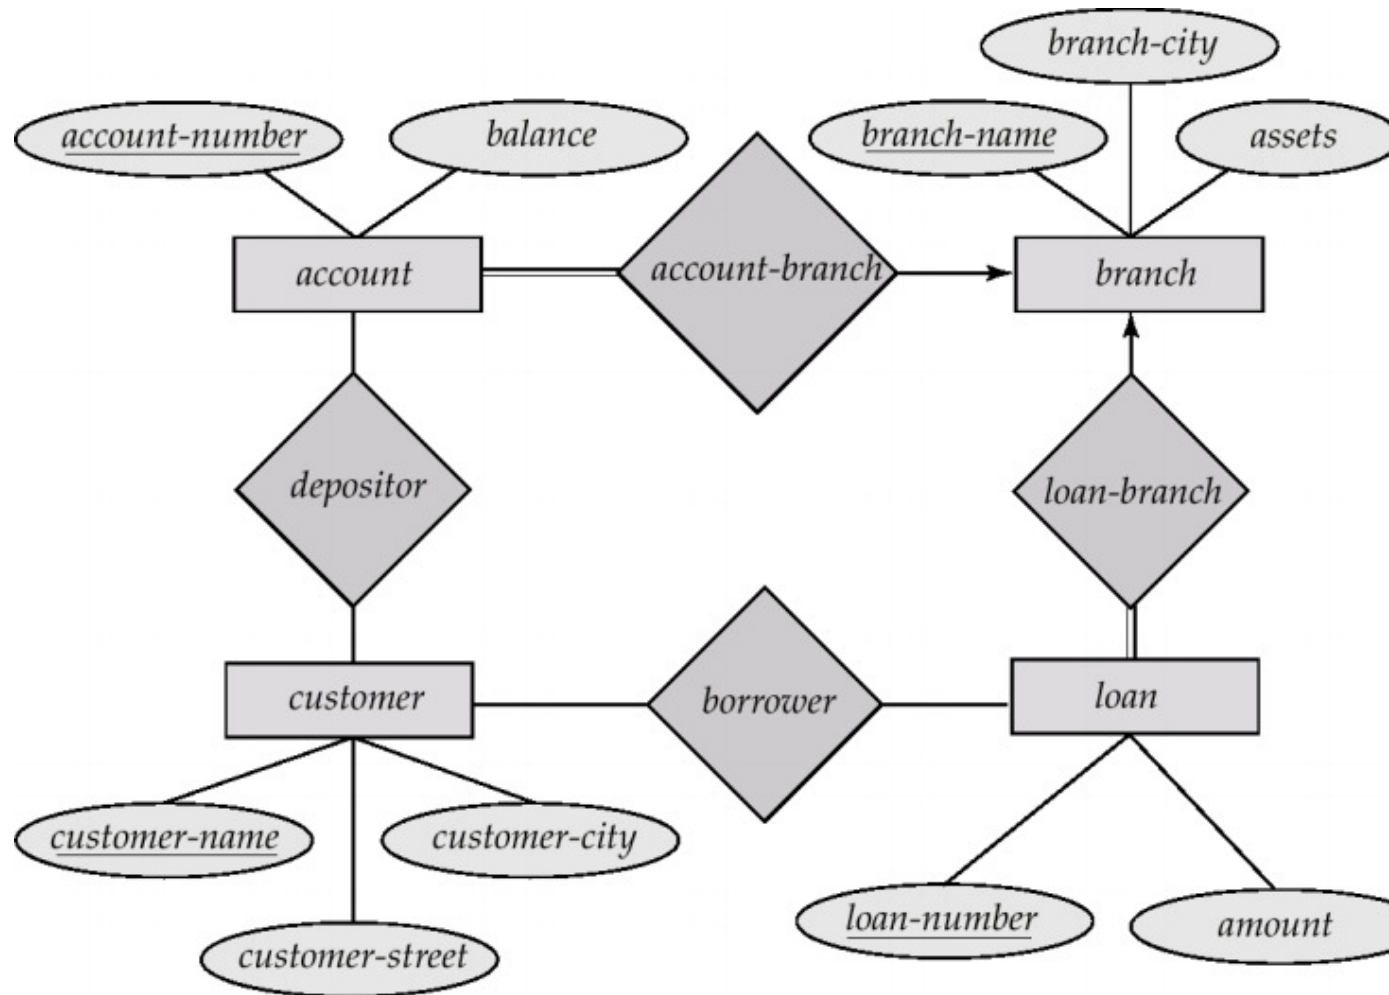
Foreign Key

- A relation schema $r1$ may include among its attributes, the primary key of another relation schema $r2$. This attribute is called a foreign key from $r1$, referencing $r2$.
- The relation $r1$ is also called the **referencing relation** of the foreign key.
- The relation $r2$ is also called the **referenced relation** of the foreign key.
- Eg. Branch-name in Account-schema is a foreign key from Account-schema referencing Branch-schema where branch-name is the primary key of Branch-schema.
- Foreign key dependencies appear as arrows from the foreign key attributes of the referencing relation to the primary key of the referenced relation.

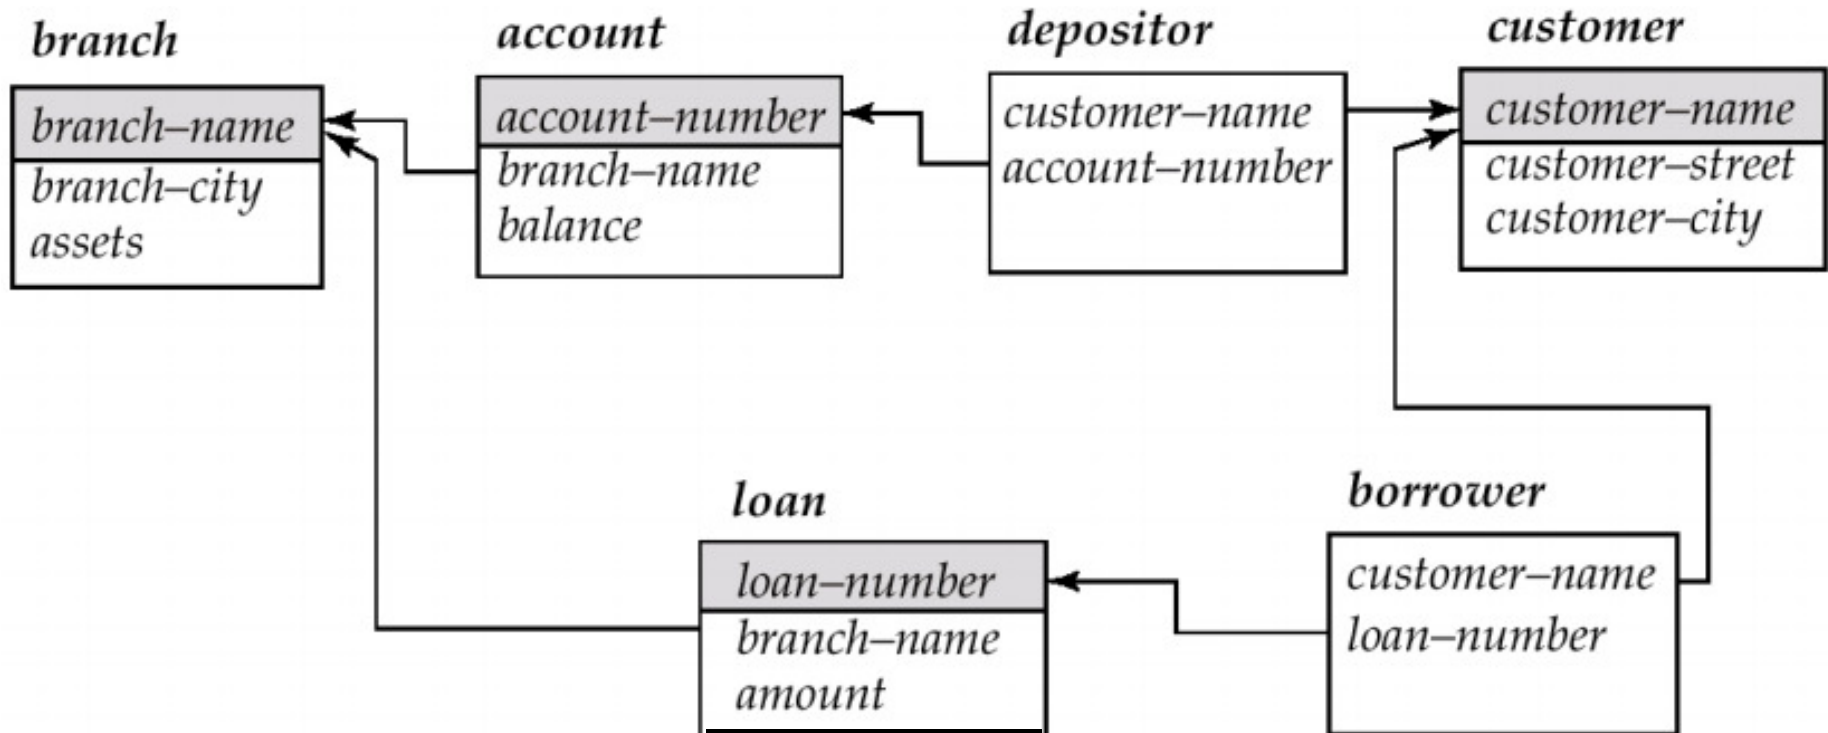
Schema diagram

- A database schema along with primary key and foreign key, can be depicted pictorially by schema diagram.
- **Boxes** represents *relations*
- **Above boxes**, are *relation names*
- **Inside boxes**, are *list of attributes*
- **Horizontal line crosses the box** for *primary key*
- **Arrows** for the *foreign key dependencies* from the foreign key attributes of the referencing relation to the primary key of the referenced relation.

E-R Diagram for the Banking Enterprise



Schema Diagram for the banking enterprise



Query Languages

- Query Language is a language in which user requests information from the database.
- Categories of languages
 - ❖ Procedural: user instruct the system to *perform a sequence of operations* on the database to compute desired result.
 - ❖ non-procedural : user describes the desired information *without giving a specific procedure* for obtaining information.
- Procedural languages:
 - ❖ Relational Algebra (procedural language)
- Non-Procedural languages:
 - ❖ Tuple Relational Calculus
 - ❖ Domain Relational Calculus
- Pure languages form underlying basis of query languages(e.g. SQL) that people use.

Relational Algebra

- Relational Algebra is a procedural query language.
- Consists of set of operations
 - that takes one or more relations as input and
 - produces a new relation as their results

- Fundamental operations

Select σ
Project Π
Union \cup
Set difference $-$
Cartesian product \times
Rename ρ

- Other operations

Set intersection \cap
Natural join \bowtie
Division \div
Assignment \leftarrow

Banking example

Consider the relation database of Bank

branch (branch-name, branch-city, assets)

customer (customer-name, customer-street, customer-city)

account (account-number, branch-name, balance)

loan (loan-number, branch-name, amount)

depositor (customer-name, account-number)

borrower (customer-name, loan-number)

Give an expression in the relational algebra to express queries.

The *Select* Operation (σ)

- Relation r

| A | B | C | D |
|----------|----------|-----|-----|
| α | α | 1 | 7 |
| α | β | 5 | 7 |
| β | β | 12 | 3 |
| β | β | 23 | 10 |

- $\sigma_{A=B \wedge D > 5}(r)$

| A | B | C | D |
|----------|----------|-----|-----|
| α | α | 1 | 7 |
| β | β | 23 | 10 |

The *Select* Operation (σ) contd..

- The select operation selects tuples that satisfy a given predicate.
- Notation: $\sigma_p(r)$
- p is called the selection predicate (selection condition).
- Predicate may consists of following clauses
 - <attribute name> <comparison operator> <attribute name>
 - <attribute name> < comparison operator > <constant value>

Comparison operators in predicate

$=, \neq, <, >, \leq, \geq$

Connectives in predicate

\wedge (And), \vee (or), \neg (Not)

The *Select* Operation (σ) contd..

E.g. Find all loans that belongs to branch name “Putalisadak”.

E.g. Find all loans that with loan amount greater than 1200.

E.g. Find all loans that belongs to branch name “Putalisadak” and loan amount greater than 1200.

| Loan-number | Branch-name | amount |
|-------------|-------------|--------|
| L-11 | Biratnagar | 900 |
| L-14 | Pokhara | 1500 |
| L-15 | Putalisadak | 1500 |
| L-16 | Putalisadak | 1300 |
| L-17 | Pokhara | 1000 |
| L-23 | Pulchowk | 2000 |
| L-93 | Birgunj | 500 |

E.g. Find all customers who have the same name as their loan officer. (*Consider a relation loan-officier which has columns customer-name, banker-name and loan-number*)

The *Select* Operation (σ) contd..

- The selection operation is commutative; that is

$$\sigma_{\langle \text{cond1} \rangle}(\sigma_{\langle \text{cond2} \rangle}(r)) = \sigma_{\langle \text{cond2} \rangle}(\sigma_{\langle \text{cond1} \rangle}(r))$$

- The cascaded SELECT operation can be combined to a single SELECT operation with a conjunctive AND(\wedge).

$$\sigma_{\langle \text{cond1} \rangle}(\sigma_{\langle \text{cond2} \rangle}(\dots \sigma_{\langle \text{condn} \rangle}(r))) = \sigma_{\langle \text{cond1} \rangle \wedge \langle \text{cond2} \rangle \wedge \dots \wedge \langle \text{condn} \rangle}(r)$$

The *Project* Operation (Π)

- Relation r

| A | B | C |
|----------|-----|-----|
| α | 10 | 1 |
| α | 20 | 1 |
| β | 30 | 1 |
| β | 40 | 2 |

- $\Pi_{A,C}(r)$

| A | C | | A | C |
|----------|-----|---|----------|-----|
| α | 1 | | α | 1 |
| α | 1 | = | β | 1 |
| β | 1 | | β | 2 |
| β | 2 | | | |

The *Project* Operation (Π) contd..

- Notation $\Pi_{A_1, A_2, \dots, A_k}(r)$
where A_1, A_2, \dots, A_k are attribute names and r is a relation.
- Selects certain columns from the table and discards other columns.
- The result is defined as the relation of k columns obtained by erasing the columns that are not listed
- Duplicate rows removed from result, since relation is a set.

Example: List all loan numbers and the amount of the loan

$\Pi_{\text{loan-number, amount}}(\text{loan})$

- Commutativity does not hold on PROJECTION

$$\Pi_{\langle \text{attribute list } 1 \rangle}(\Pi_{\langle \text{attribute list } 2 \rangle}(r)) = \Pi_{\langle \text{attribute list } 1 \rangle}(r)$$

Combination of relational operators

- The result of a relational algebra operation is itself a relation.
- Can build a relational-algebra expression using multiple operations
- E.g. Find all the customer who lives in “Harrison”.

| <i>customer-name</i> | <i>customer-street</i> | <i>customer-city</i> |
|----------------------|------------------------|----------------------|
| Adams | Spring | Pittsfield |
| Brooks | Senator | Brooklyn |
| Curry | North | Rye |
| Glenn | Sand Hill | Woodside |
| Green | Walnut | Stamford |
| Hayes | Main | Harrison |
| Johnson | Alma | Palo Alto |
| Jones | Main | Harrison |
| Lindsay | Park | Pittsfield |
| Smith | North | Rye |
| Turner | Putnam | Stamford |
| Williams | Nassau | Princeton |

$\Pi_{customer-name} (\sigma_{customer-city="Harrison"} (r))$

The *Union* Operation (\cup)

- Relation r, s

| A | B |
|----------|-----|
| α | 1 |
| α | 2 |
| β | 1 |

r

| A | B |
|----------|-----|
| α | 2 |
| β | 3 |

s

- $r \cup s$

| A | B |
|----------|-----|
| α | 1 |
| α | 2 |
| β | 1 |
| β | 3 |

The *Union* Operation (\cup) contd..

- Notation: $r \cup s$
- The result of this operation is a relation that **includes all tuples that are either in r or in s or in both.**
- Duplicate tuples are eliminated.
- Union must be taken between compatible relations, that means, for $r \cup s$ to be valid, the two conditions must hold.
 1. The ***relations r and s*** must have the ***same arity***. That is, they must have ***same number of attributes***.
 2. The domains of the **i^{th} attribute of r** and the **i^{th} attribute of s** must be the **same, for all i .**
- Set union is commutative, that is

$$r \cup s = s \cup r$$

- E.g. to find all customers with either an account or a loan

$$\Pi_{\text{customer-name}}(\text{borrower}) \cup \Pi_{\text{customer-name}}(\text{depositor})$$

Set Difference (—)

- Relations r, s :

| A | B |
|----------|-----|
| α | 1 |
| α | 2 |
| β | 1 |

r

| A | B |
|----------|-----|
| α | 2 |
| β | 3 |

s

- $r - s$

| A | B |
|----------|-----|
| α | 1 |
| β | 1 |

Set Difference (—) contd..

- Notation: $r - s$
- The result of this operation is a relation that includes all tuples that are in r but not in s .
- Set differences must be taken between compatible relations.
 1. The *relations r and s* must have the *same arity*. That is, they must have *same number of attributes*.
 2. The domains of the i^{th} attribute of r and the i^{th} attribute of s must be the same, for all i .
- Set difference is not commutative, that is
$$r - s \neq s - r$$

Example:

E.g. to find all customers who have an account but not a loan

$$\Pi_{\text{customer-name}}(\text{depositor}) - \Pi_{\text{customer-name}}(\text{borrower})$$

Cartesian Product (\times)

- Relations r, s :

| A | B |
|----------|-----|
| α | 1 |
| β | 2 |

r

| C | D | E |
|----------|-----|-----|
| α | 10 | a |
| β | 10 | a |
| β | 20 | b |
| γ | 10 | b |

s

- $r \times s$

| A | B | C | D | E |
|----------|-----|----------|-----|-----|
| α | 1 | α | 10 | a |
| α | 1 | β | 10 | a |
| α | 1 | β | 20 | b |
| α | 1 | γ | 10 | b |
| β | 2 | α | 10 | a |
| β | 2 | β | 10 | a |
| β | 2 | β | 20 | b |
| β | 2 | γ | 10 | b |

Cartesian Product (\times) contd..

- Notation $r1 \times r2$
- Also called *cross product*
- This operation is used to *combine tuples from any two relations.*
- In general, if *relations $r1(R1)$ and $r2(R2)$, then $r1 \times r2$ is a relation*
 1. *whose schema is the concatenation of $R1$ and $R2$.*
 2. *which contains all the tuples t for which there is a tuple $t1$ in $r1$ and a tuple $t2$ in $r2$ for which $t[R1]=t1[R1]$ and $t[R2]=t2[R2]$*
- *If we have $n1$ tuples $r1$ and $n2$ tuples in $r2$, then there are $n1 * n2$ ways of choosing pair of tuples, so there is $n1 * n2$ tuples in r , where $r=r1 \times r2$.*

Cartesian Product (\times) contd..

- If attributes of relation schemas R1 and R2 are not disjoint, then,
 - such attributes are used by attaching to an attribute the name of the relation from which the attribute belongs to.

OR

- the renaming must be done.
- *Example:*
borrower (customer-name, loan-number)
loan (loan-number, branch-name, amount)

The relation schema for $r = \text{borrower} \times \text{loan}$ is

(borrower.customer-name, borrower.loan-number, loan.loan-number, loan.branch-name, loan.amount)

This can also be written without any ambiguity

(customer-name, borrower.loan-number, loan.loan-number, branch-name, amount)

Cartesian Product (x) Example

Find names of all customers who have a loan at Putalisadak branch.

- Step 1: we need information from both customer and loan tables

$\text{Borrower} \times \text{loan}$

- Step 2: Find tuples which belongs to branch “Putalisadak”

$\sigma_{\text{branch.name}=\text{“Putalisadak”}}(\text{Borrower} \times \text{loan})$

- Step 3: Find tuples that pertain to customers who have loan at Putalisadak branch

$\sigma_{\text{borrower.loan-number}=\text{loan.loan-number}}($

$\sigma_{\text{branch.name}=\text{“Putalisadak”}}(\text{Borrower} \times \text{loan}))$

Step 4: We only want to show customer-name

$\pi_{\text{customer-name}}($

$\sigma_{\text{borrower.loan-number}=\text{loan.loan-number}}($

$\sigma_{\text{branch.name}=\text{“Putalisadak”}}(\text{Borrower} \times \text{loan}))$

Example

Find names of all customers who have a loan at Putalisadak branch.

| customer-name | loan-number |
|---------------|-------------|
| Shankar | L-15 |
| John | L-16 |
| Ashok | L-93 |
| Anil | L-14 |

Table *borrower*

| loan-numberc | branch-name | amount |
|--------------|-------------|--------|
| L-14 | Pokhara | 1500 |
| L-15 | Putalisadak | 1500 |
| L-16 | Putalisadak | 1300 |
| L-93 | Birgunj | 500 |

Table *Loan*

Borrower x loan

| customer-name | borrower.loan -number | loan.loan -number | branch-name | amount |
|---------------|--------------------------|----------------------|-------------|--------|
| Shankar | L-15 | L-14 | Pokhara | 1500 |
| Shankar | L-15 | L-15 | Putalisadak | 1500 |
| Shankar | L-15 | L-16 | Putalisadak | 1300 |
| Shankar | L-15 | L-93 | Birgunj | 500 |
| John | L-16 | L-14 | Pokhara | 1500 |
| John | L-16 | L-15 | Putalisadak | 1500 |
| John | L-16 | L-16 | Putalisadak | 1300 |
| John | L-16 | L-93 | Birgunj | 500 |
| Ashok | L-93 | L-14 | Pokhara | 1500 |
| Ashok | L-93 | L-15 | Putalisadak | 1500 |
| Ashok | L-93 | L-16 | Putalisadak | 1300 |
| Ashok | L-93 | L-93 | Birgunj | 500 |
| Anil | L-14 | L-14 | Pokhara | 1500 |
| Anil | L-14 | L-15 | Putalisadak | 1500 |
| Anil | L-14 | L-16 | Putalisadak | 1300 |
| Anil | L-14 | L-93 | Birgunj | 500 |

$\sigma_{\text{branch.name="Putalisadak"}}(\text{Borrower} \times \text{loan})$

| customer-name | borrower.loan -number | loan.loan -number | branch-name | amount |
|--------------------|--------------------------|----------------------|--------------------|-----------------|
| Shankar | L-15 | L-14 | Pokhara | 1500 |
| Shankar | L-15 | L-15 | Putalisadak | 1500 |
| Shankar | L-15 | L-16 | Putalisadak | 1300 |
| Shankar | L-15 | L-93 | Birgunj | 500 |
| John | L-16 | L-14 | Pokhara | 1500 |
| John | L-16 | L-15 | Putalisadak | 1500 |
| John | L-16 | L-16 | Putalisadak | 1300 |
| John | L-16 | L-93 | Birgunj | 500 |
| Ashok | L-93 | L-14 | Pokhara | 1500 |
| Ashok | L-93 | L-15 | Putalisadak | 1500 |
| Ashok | L-93 | L-16 | Putalisadak | 1300 |
| Ashok | L-93 | L-93 | Birgunj | 500 |
| Anil | L-14 | L-14 | Pokhara | 1500 |
| Anil | L-14 | L-15 | Putalisadak | 1500 |
| Anil | L-14 | L-16 | Putalisadak | 1300 |
| Anil | L-14 | L-93 | Birgunj | 500 |

Step 3

| customer-name | borrower.loan -number | loan.loan -number | branch-name | amount |
|---------------|--------------------------|----------------------|-------------|--------|
| Shankar | L-15 | L-15 | Putalisadak | 1500 |
| Shankar | L-15 | L-16 | Putalisadak | 1300 |
| John | L-16 | L-15 | Putalisadak | 1500 |
| John | L-16 | L-16 | Putalisadak | 1300 |
| Ashok | L-93 | L-15 | Putalisadak | 1500 |
| Ashok | L-93 | L-16 | Putalisadak | 1300 |
| Anil | L-14 | L-15 | Putalisadak | 1500 |
| Anil | L-14 | L-16 | Putalisadak | 1300 |

$\sigma_{\text{borrower.loan-number}=\text{loan.loan-number}}$
 $\sigma_{\text{branch.name}=\text{"Putalisadak"}}(\text{Borrower x loan})$

Step 4

| customer-name | borrower loan -number | loan.loan -number | branch-name | amount |
|---------------|--------------------------|----------------------|-------------|--------|
| Shankar | L-15 | L-15 | Putalisadak | 1500 |
| John | L-16 | L-16 | Putalisadak | 1300 |

$\Pi_{\text{customer-name}}(\sigma_{\text{borrower.loan-number}=\text{loan.loan-number}}(\sigma_{\text{branch.name}=\text{"Putalisadak"}}(\text{Borrower x loan})))$

Rename (ρ)

- The symbol ρ (rho) denotes the rename operator.
- Allows us to name the results of relational algebra expression E
 $\rho_x(E)$ returns the result of expression under name x .
- Allows rename either the relation name, attribute names or both.
- Thus, if a relation algebra expression E has arity n (i.e. n columns), then

$\rho_{x(A_1, A_2, \dots, A_n)}(E)$ returns the result of expression under name x and attributes renamed to A_1, A_2, \dots, A_n

$\rho_{x(A_1, A_2, \dots, A_n)}(r)$ renames both relation and its attributes.

$\rho_x(r)$ renames relation only.

$\rho_{(A_1, A_2, \dots, A_n)}(r)$ renames attributes only.

If the attributes of r are C_1, C_2, \dots, C_n in that order, then each C_i is renamed as A_i .

Rename (ρ) contd..

- Find a largest account balance in the bank

| <i>account-number</i> | <i>branch-name</i> | <i>balance</i> |
|-----------------------|--------------------|----------------|
| A-101 | Downtown | 500 |
| A-215 | Mianus | 700 |
| A-102 | Perryridge | 400 |
| A-305 | Round Hill | 350 |
| A-201 | Brighton | 900 |
| A-222 | Redwood | 700 |
| A-217 | Brighton | 750 |

Rename (ρ) contd..

Step 1. gives those balances except the largest balance

$$\Pi_{account.balance}(\sigma_{account.balance < d.balance}(\text{account} \times \rho_d(\text{account})))$$

Step 2. applying set difference gives the largest balance in a bank

$$\Pi_{account.balance}(\text{account}) - \Pi_{account.balance}(\sigma_{account.balance < d.balance}(\text{account} \times \rho_d(\text{account})))$$

Rename (ρ) contd..

Find the names of all customers who live on the same street and city as smith.

| <i>customer-name</i> | <i>customer-street</i> | <i>customer-city</i> |
|----------------------|------------------------|----------------------|
| Adams | Spring | Pittsfield |
| Brooks | Senator | Brooklyn |
| Curry | North | Rye |
| Glenn | Sand Hill | Woodside |
| Green | Walnut | Stamford |
| Hayes | Main | Harrison |
| Johnson | Alma | Palo Alto |
| Jones | Main | Harrison |
| Lindsay | Park | Pittsfield |
| Smith | North | Rye |
| Turner | Putnam | Stamford |
| Williams | Nassau | Princeton |

Rename (ρ) contd..

Step 1. get Smith's street-name and city by an expression

$\Pi_{customer-street, customer-city}(\sigma_{customer-name="Smith"}(customer))$

Step 2. find other customers who lives in Smith's street and city

$\Pi_{account.balance}(\sigma_{customer.customer-street=smith-addr.street}$
 $\wedge customer.customer-city=smith-addr.city$
 $(customer \times \rho_{smith-addr(street, city)}$
 $(\Pi_{customer-street, customer-city}(\sigma_{customer-name="Smith"}(customer))))$

Formal Definition of the Relational Algebra

- A basic expression in the relational algebra consists of either one of the following:
 - ❖ A relation in the database
 - ❖ A constant relation
- A constant relation is written by listing its tuples within { }.
e.g. account = { (A-101, Biratnagar, 500) (A-215, Pokhara, 700) }

A general expression in relational algebra composed of smaller expressions. Let $E1$ and $E2$ are relational algebra expressions. Then, relation algebra expressions are

a. $E1 \cup E2$

b. $E1 - E2$

c. $E1 \times E2$

d. $\sigma_P(E)$, where P is a predicate on attributes in $E1$

e. $\Pi_S(E1)$, where S is a list consisting of some attributes

f. $\rho_x(E1)$, where x is the new name for the result of $E1$

Additional Operations

- The fundamental operations of relational algebra are sufficient to express any query.
- Certain common queries are lengthy to express only by using fundamental operations
- Additional operations that do not add any power to the relational algebra, but that simplify common queries.
- Additional Algebra Operations
 - Set intersection \cap
 - Natural join \bowtie
 - Division \div
 - Assignment \leftarrow

Set-Intersection Operation \cap

- *Relation r, s :*

| A | B |
|----------|---|
| α | 1 |
| α | 2 |
| β | 1 |

r

| A | B |
|----------|---|
| α | 2 |
| β | 3 |

s

- $r \cap s$

| A | B |
|----------|---|
| α | 2 |

Set-Intersection Operation \cap

- Notation: $r \cap s$
- The result of this operation, denoted by $r \cap s$, is a relation that includes all tuples that are in both r and s .
- For intersection, r and s must be compatible relations.

E.g. Find all customers who have both a loan and an account.

$$\Pi_{customer-name}(borrower) \cap \Pi_{customer-name}(depositor)$$

Note: The equivalent relational algebra expressions

$$r \cap s = r - (r - s)$$

Natural-Join Operation \bowtie

- Notation: $r \bowtie s$
- Consider two relations $r(R)$ and $s(S)$ where R and S are schemas.
 - ❖ *The result is a relation on schema $R \cup S$ which is obtained by considering each pair of tuples t_r from r and t_s from s .*
 - ❖ *If t_r and t_s have the same value on each of the attributes in $R \cap S$, a tuple t is added to the result. Where*
 - *t has the same value as t_r on r .*
 - *t has the same value as t_s on s .*

Example:

$R=(A, B, C, D)$

$S=(E, B, D)$

Result schema=(A, B, C, D, E)

$r \bowtie s$ is defined as

$$\Pi_{r.A, r.B, r.C, r.D, s.E} (\sigma_{r.B=s.B \wedge r.D=s.D} (r \times s))$$

Natural-Join Operation \bowtie

- Relation r, s

| A | B | C | D |
|----------|-----|----------|-----|
| α | 1 | α | a |
| β | 2 | γ | a |
| γ | 4 | β | b |
| α | 1 | γ | a |
| δ | 2 | β | b |

r

| B | D | E |
|-----|-----|------------|
| 1 | a | α |
| 3 | a | β |
| 1 | a | γ |
| 2 | b | δ |
| 3 | b | ϵ |

s

- $r \bowtie s$

| A | B | C | D | E |
|----------|-----|----------|-----|----------|
| α | 1 | α | a | α |
| α | 1 | α | a | γ |
| α | 1 | γ | a | α |
| α | 1 | γ | a | γ |
| δ | 2 | β | b | δ |

Natural-Join Operation

- Consider two relations $r(R)$ and $s(S)$
- The natural join of r and s , denoted by $r \bowtie s$, is a relation on schema $R \cup S$ formally defined as follows

$$r \bowtie s = \Pi_{R \cup S} (\sigma_{r.A_1=s.A_1 \wedge r.A_2=s.A_2 \wedge \dots \wedge r.A_n=s.A_n}(r \times s))$$

Where $R \cap S = \{A_1, A_2, \dots, A_n\}$

Table *account*

| <i>account-number</i> | <i>branch-name</i> | <i>balance</i> |
|-----------------------|--------------------|----------------|
| A-101 | Downtown | 500 |
| A-102 | Perryridge | 400 |
| A-201 | Brighton | 900 |
| A-215 | Mianus | 700 |
| A-217 | Brighton | 750 |
| A-222 | Redwood | 700 |
| A-305 | Round Hill | 350 |

Table *depositor*

| <i>customer-name</i> | <i>account-number</i> |
|----------------------|-----------------------|
| Hayes | A-102 |
| Johnson | A-101 |
| Johnson | A-201 |
| Jones | A-217 |
| Lindsay | A-222 |
| Smith | A-215 |
| Turner | A-305 |

Table *customer*

| <i>customer-name</i> | <i>customer-street</i> | <i>customer-city</i> |
|----------------------|------------------------|----------------------|
| Adams | Spring | Pittsfield |
| Brooks | Senator | Brooklyn |
| Curry | North | Rye |
| Glenn | Sand Hill | Woodside |
| Green | Walnut | Stamford |
| Hayes | Main | Harrison |
| Johnson | Alma | Palo Alto |
| Jones | Main | Harrison |
| Lindsay | Park | Pittsfield |
| Smith | North | Rye |
| Turner | Putnam | Stamford |
| Williams | Nassau | Princeton |

1. Find the names of all branches with customers who have an account in the bank and who live in harrison.
2. Find all the customer who have both a loan and an account at the bank.

theta-Join Operation

The theta join is an extension to the natural join operation which allows to combine a selection and a cartesian product.

Consider relations $r(R)$ and $s(S)$, θ be the predicate on attributes in the schema $R \cup S$. Then theta join operation $r \bowtie_{\theta} s$ is defined as

$$r \bowtie_{\theta} s = \sigma_{\theta} (r \times s)$$

E.g.

Department $\bowtie_{\text{MGRSSN=SSN}}$ Employee

Division Operation \div

- Division operation denoted by \div is suited to queries that includes the phrase “for all”.
- Formally, let two relations $r(R)$ and $s(S)$ and let $S \subseteq R$. That is, every attributes of schema S is also in R .

The relation $r \div s$ is a relation on schema $R - S$.

A tuple t is in $r \div s$ if and only if both of two conditions holds:

- 1. t is in $\Pi_{R-S}(r)$*
- 2. For every tuple t_s in s , there is a tuple t_r in r satisfying both of following:*
 - a. $t_r[S] = t_s[S]$*
 - b. $t_r[R-S] = t$*

Division Operation \div

Example 1

- Relation r, s

| A | B |
|------------|---|
| α | 1 |
| α | 2 |
| α | 3 |
| β | 1 |
| γ | 1 |
| δ | 1 |
| δ | 3 |
| δ | 4 |
| ϵ | 6 |
| ϵ | 1 |
| β | 2 |

r

| B |
|-----|
| 1 |
| 2 |
| s |

- $r \div s$

| A |
|----------|
| α |
| β |

- t is in $\Pi_{R-S}(r)$
- For every tuple t_s in s , there is a tuple t_r in r satisfying both of following:
 - $t_r[S] = t_s[S]$
 - $t_r[R-S] = t$

Division Operation \div

Example 2

- Relation r, s

| A | B | C | D | E |
|----------|---|----------|---|---|
| α | a | α | a | 1 |
| α | a | γ | a | 1 |
| α | a | γ | b | 1 |
| β | a | γ | a | 1 |
| β | a | γ | b | 3 |
| γ | a | γ | a | 1 |
| γ | a | γ | b | 1 |
| γ | a | β | b | 1 |

r

| D | E |
|---|---|
| a | 1 |
| b | 1 |

s

- $r \div s$

| A | B | C |
|----------|---|----------|
| α | a | γ |
| γ | a | γ |

Example

- Find all the customers who have an account at all the branches located in Brooklyn.

| Branch-name | Branch-city | assets |
|-------------|-------------|---------|
| Brighton | Brooklyn | 7100000 |
| Downtown | Brooklyn | 9000000 |
| Mianus | Horseneck | 400000 |
| .. | | |
| | | |

Table *branch*

Assignment Operators \leftarrow

- The assignment operation (\leftarrow) provides a convenient way to express complex queries, write query as a sequential program consisting of a series of assignments followed by an expression whose value is displayed as a result of the query.
- Assignment must always be made to a temporary relation variable.
- Example $r \div s$ can be written as
$$\text{temp1} \leftarrow \Pi_{R-S}(r)$$
$$\text{temp2} \leftarrow \Pi_{R-S}(\text{temp1} \times s) - \Pi_{R-S, S}(r)$$
$$\text{Result} \leftarrow \text{temp1} - \text{temp2}$$
- The result to the right of the \leftarrow is assigned to the relation variable on the left of the \leftarrow .

Extended Relational-Algebra-Operations

- Generalized Projection
- Outer Join
- Aggregate Functions

Generalized Projection

- Extends the projection operation by allowing arithmetic functions to be used in the projection list.

$$\Pi_{F_1, F_2, \dots, F_n}(E)$$

- E is any relational-algebra expression.
- Each of F_1, F_2, \dots, F_n is an arithmetic expression involving constants and attributes in the schema of E .
- An arithmetic expression may be simply an attributes or a constant.

Generalized Projection contd..

Table *employee*

| ssn | name | salary | tax |
|--------|-------|--------|------|
| 111221 | John | 10000 | 0.05 |
| 511249 | Smith | 20000 | 0.1 |

Write a query to
find the employee net earning.

Table *result of the expression*

| ssn | name | net-salary |
|--------|-------|------------|
| 111221 | John | |
| 511249 | Smith | |

Aggregate functions

- Aggregation function takes a collection of values and returns single value as a result.
 - avg**: average value
 - min**: minimum value
 - max**: maximum value
 - sum**: sum of values
 - count**: number of values
- General form of the aggregation operation

$$G_1, G_2, \dots, G_n \text{ } g \text{ } F_1(A_1), F_2(A_2), \dots, F_n(A_n) (E)$$

- E is any relational-algebra expression
- G_1, G_2, \dots, G_n is list of attributes on which to group (can be empty).
- Each F_i is an aggregate function.
- Each A_i is an attribute name.

Aggregate functions- Example

- Relation r

| <i>A</i> | <i>B</i> | <i>C</i> |
|----------|----------|----------|
| α | α | 7 |
| α | β | 7 |
| β | β | 3 |
| β | β | 10 |

| |
|--------------|
| <i>sum-C</i> |
| 27 |

Aggregate functions- Example

Table *account*

| <i>branch-name</i> | <i>account-number</i> | <i>balance</i> |
|--------------------|-----------------------|----------------|
| Perryridge | A-102 | 400 |
| Perryridge | A-201 | 900 |
| Brighton | A-217 | 750 |
| Brighton | A-215 | 750 |
| Redwood | A-222 | 700 |

Write an algebra expression
to find sum of all balance that belongs to each branch.

| <i>branch-name</i> | <i>balance</i> |
|--------------------|----------------|
| Perryridge | 1300 |
| Brighton | 1500 |
| Redwood | 700 |

Aggregate functions contd..

- Result of aggregation does not have a name
- Can use rename operation to give it a name.
- For convenience, we permit renaming as part of aggregate operation

branch-name ***g*** *sum(balance)* ***as*** *sum-balance* (*account*)

Aggregate functions contd..

| A | B | value |
|----|----|-------|
| a1 | a1 | 11 |
| a1 | a1 | 100 |
| b1 | a1 | 22 |
| b1 | a1 | 30 |
| a1 | a1 | 6 |
| b1 | a1 | 50 |

table *r*

Write an algebra expression to get following result relation from table *r*.

Result a.

| B | mininum | maximum |
|----|---------|---------|
| a1 | 11 | 100 |
| a1 | 22 | 50 |

Result b.

| A | B | mininum | maximum |
|----|----|---------|---------|
| a1 | a1 | 11 | 100 |
| b1 | a1 | 6 | 50 |

Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values
 - ❖ *null* signifies that the value is unknown or does not exist
 - ❖ All comparison between nulls are false.

Inner Join

Table *loan*

| <i>loan-number</i> | <i>branch-name</i> | <i>amount</i> |
|--------------------|--------------------|---------------|
| L-170 | Downtown | 3000 |
| L-230 | Redwood | 4000 |
| L-260 | Perryridge | 1700 |

Table *customer*

| <i>customer-name</i> | <i>loan-number</i> |
|----------------------|--------------------|
| Jones | L-170 |
| Smith | L-230 |
| Hayes | L-155 |

Inner join:

loan ⋈ *customer*

| <i>loan-number</i> | <i>branch-name</i> | <i>amount</i> | <i>customer-name</i> |
|--------------------|--------------------|---------------|----------------------|
| L-170 | Downtown | 3000 | Jones |
| L-230 | Redwood | 4000 | Smith |

Left OuterJoin

The left outer join \bowtie of any two relations is

- the result of natural join plus
- all tuples in the left relation that did not match with any tuples in the right relations, pads the tuples with null values for all other attributes

Table *loan*

| <i>loan-number</i> | <i>branch-name</i> | <i>amount</i> |
|--------------------|--------------------|---------------|
| L-170 | Downtown | 3000 |
| L-230 | Redwood | 4000 |
| L-260 | Perryridge | 1700 |

Table *customer*

| <i>customer-name</i> | <i>loan-number</i> |
|----------------------|--------------------|
| Jones | L-170 |
| Smith | L-230 |
| Hayes | L-155 |

Left outer join:

loan \bowtie *customer*

| <i>loan-number</i> | <i>branch-name</i> | <i>amount</i> | <i>customer-name</i> |
|--------------------|--------------------|---------------|----------------------|
| L-170 | Downtown | 3000 | Jones |
| L-230 | Redwood | 4000 | Smith |
| L-260 | Perryridge | 1700 | <i>null</i> |

Right Outer Join

The right outer join \bowtie on any two relations is

- the result of natural join plus
- all tuples in the right relation that did not match with any tuples in the left relations, pads the tuples with null values for all other attributes

Table *loan*

| <i>loan-number</i> | <i>branch-name</i> | <i>amount</i> |
|--------------------|--------------------|---------------|
| L-170 | Downtown | 3000 |
| L-230 | Redwood | 4000 |
| L-260 | Perryridge | 1700 |

Table *customer*

| <i>customer-name</i> | <i>loan-number</i> |
|----------------------|--------------------|
| Jones | L-170 |
| Smith | L-230 |
| Hayes | L-155 |

Right outer join:

loan \bowtie *customer*

| <i>loan-number</i> | <i>branch-name</i> | <i>amount</i> | <i>customer-name</i> |
|--------------------|--------------------|---------------|----------------------|
| L-170 | Downtown | 3000 | Jones |
| L-230 | Redwood | 4000 | Smith |
| L-155 | <i>null</i> | <i>null</i> | Hayes |

Full Outer Join

Table *loan*

| <i>loan-number</i> | <i>branch-name</i> | <i>amount</i> |
|--------------------|--------------------|---------------|
| L-170 | Downtown | 3000 |
| L-230 | Redwood | 4000 |
| L-260 | Perryridge | 1700 |

Table *customer*

| <i>customer-name</i> | <i>loan-number</i> |
|----------------------|--------------------|
| Jones | L-170 |
| Smith | L-230 |
| Hayes | L-155 |

Right outer join:

$\text{loan} \bowtie \text{customer}$

| <i>loan-number</i> | <i>branch-name</i> | <i>amount</i> | <i>customer-name</i> |
|--------------------|--------------------|---------------|----------------------|
| L-170 | Downtown | 3000 | Jones |
| L-230 | Redwood | 4000 | Smith |
| L-260 | Perryridge | 1700 | <i>null</i> |
| L-155 | <i>null</i> | <i>null</i> | Hayes |

Modification of database

- The content of the database may be modified using the following operations:
 - ❖ Deletion
 - ❖ Insertion
 - ❖ Updating
- Using assignment operation \leftarrow database modifications are done

Deletion

- A delete request is expressed similarly to a query, except instead of displaying tuples to the user, the selected tuples are removed from the database.
- Can delete only whole tuples; cannot delete values on only particular attribute.
- A deletion is expressed in relational algebra by:

$$r \leftarrow r - E$$

where r is a relational and E is a relation algebra expression.

Table *account*

| <i>account-number</i> | <i>branch-name</i> | <i>balance</i> |
|-----------------------|--------------------|----------------|
| A-101 | Downtown | 500 |
| A-215 | Mianus | 700 |
| A-102 | Perryridge | 400 |
| A-305 | Round Hill | 350 |
| A-201 | Brighton | 900 |
| A-222 | Redwood | 700 |
| A-217 | Brighton | 750 |

Table *depositor*

| <i>customer-name</i> | <i>account-number</i> |
|----------------------|-----------------------|
| Hayes | A-102 |
| Johnson | A-101 |
| Johnson | A-201 |
| Jones | A-217 |
| Lindsay | A-222 |
| Smith | A-215 |
| Turner | A-305 |

- Delete all of Smith's account records from depositor
- Delete all loan records with amount in the range of 0 to 50
- Delete all accounts at branches located in Brooklyn.

| <i>branch-name</i> | <i>branch-city</i> | <i>assets</i> |
|--------------------|--------------------|---------------|
| Brighton | Brooklyn | 7100000 |
| Downtown | Brooklyn | 9000000 |
| Mianus | Horseneck | 400000 |

Deletion contd..

- Delete all of Smith's account records from depositor
$$\text{depositor} \leftarrow \text{depositor} - \sigma_{\text{customer-name}=\text{"smith"}}(\text{depositor})$$
- Delete all loan records with amount in the range of 0 to 50
$$\text{loan} \leftarrow \text{loan} - \sigma_{\text{amount} \geq 0 \text{ and } \text{amount} \leq 50}(\text{loan})$$
- Delete all accounts at branches located in Brooklyn.
$$r1 \leftarrow \sigma_{\text{branch-city}=\text{"Brooklyn"}}(\text{account} \bowtie \text{branch})$$
$$r2 \leftarrow \Pi_{\text{branch-name, account-number, balance}}(r1)$$
$$\text{account} \leftarrow \text{account} - r1$$

Insertion

- To insert data into a relation, we either
 - ❖ specify a tuple to be inserted
 - ❖ write a query whose result is a set of tuples to be inserted
- in relational algebra, an insertion is expressed by:
$$r \leftarrow r \cup E$$

r is a relational and E is a relation algebra expression.
- The insertion of a single tuple is expressed by letting E be a constant relation containing one tuple.

Insertion contd..

- Insert information in the database specifying that Smith has \$1200 in account A-973 at the Perryridge branch.
- Provide as a gift for all loan customers in the Perryridge branch, a \$200 savings account. Let the loan number serve as the account number for the new savings account.

Table *account*

| <i>account-number</i> | <i>branch-name</i> | <i>balance</i> |
|-----------------------|--------------------|----------------|
| A-101 | Downtown | 500 |
| A-215 | Mianus | 700 |
| A-102 | Perryridge | 400 |

Table *depositor*

| <i>customer-name</i> | <i>account-number</i> |
|----------------------|-----------------------|
| Hayes | A-102 |
| Johnson | A-101 |
| Johnson | A-201 |

Insertion contd..

- Insert information in the database specifying that Smith has \$1200 in account A-973 at the Perryridge branch.

$\text{account} \leftarrow \text{account} \cup \{(A-973, \text{"Perryridge"}, 1200)\}$
 $\text{depositor} \leftarrow \text{depositor} \cup \{(\text{"smith"}, A-973)\}$

- Provide as a gift for all loan customers in the Perryridge branch, a \$200 savings account. Let the loan number serve as the account number for the new savings account.

$\text{temp1} \leftarrow \sigma_{\text{branch-name}=\text{"Perryridge"}} (\text{borrower} \bowtie \text{loan})$

$\text{temp2} \leftarrow \Pi_{\text{loan-number}, \text{branch-name}}(\text{temp1})$

$\text{account} \leftarrow \text{account} \cup (\text{temp2} \times \{200\})$

$\text{depositor} \leftarrow \text{depositor} \cup \Pi_{\text{customer-name}, \text{loan-number}}(\text{temp1})$

Updating

- A mechanism to change a value in a tuple without changing all values in the tuple
- Use the generalized projection operator to do this task

$$r \leftarrow \Pi_{F_1, F_2, \dots, F_n}(r)$$

Each F_i is either the i^{th} attribute of r , if the i^{th} attribute is not updated

OR

if the attribute is to be updated, F_i is an expression involving only constants and the attributes of F_i , which gives the new value for the attribute

Updating contd..

E.g. 1. Interest payments are being changed and that all balances are to be increased by 5 percent.

E.g. 2. accounts with balances over \$10,000 receive 6 percent interest, whereas all other receive 5 percent.

Updating contd..

E.g. 1. Interest payments are being changed and that all balances are to be increased by 5 percent.

$account \leftarrow \Pi_{account\text{-}number, branch\text{-}name, balance * 1.05} (account)$

E.g. 2. accounts with balances over \$10,000 receive 6 percent interest, whereas all other receive 5 percent.

$account \leftarrow \Pi_{account\text{-}number, branch\text{-}name, balance * 1.06} (\sigma_{balance > 10000} (account))$

U

$\Pi_{account\text{-}number, branch\text{-}name, balance * 1.05} (\sigma_{balance \leq 10000} (account))$

End of chapter 3